

Sleeping Machines — investor pitch

A learning architecture and computing substrate for intelligence everywhere.

We are building AI that computes through time and sparse events, learns which work to perform, and maintains persistent context. The ambition is capable, adaptive intelligence from datacenters to mobile devices and robots, with models, execution software and eventually clockless hardware designed together.

Our core combines learned temporal races, addressed memory, separate keys and values, deep content-bearing messages and counterfactual credit to routes that did not win. This creates a path to more useful capacity than selected work, and to learning near the point of use. The complete construction and its quality/resource consequences are the proposed differentiation.

The research already has tangible positive results. On the completed single-seed 10M-character, one-pass text8 comparison, a small credited native model beats the saved Transformer control with about 15.4 times less estimated fitting work. A larger model slightly beats the saved LSTM's quality but uses more work. A separate causal CPU pilot improves predictions through full-backbone online adaptation. These are early scoped results, not modern frontier or hardware claims. The current report preserves their controls and limitations.

The first commercial target is useful quality at lower total datacenter cost. The next steps are actual-trained sparse-backend admission, measured serving, modern replication and a defined customer deployment. Software/model licensing can precede accelerator IP and custom silicon. Mobile and robotics extend the platform toward always-available contextual AI and adaptive machines.

The upside is ownership of a widely used AI infrastructure layer: models and runtime, licensed compute IP, and potentially a family of learning-capable event processors. Repeatable datacenter deployments or OEM design wins can support large businesses; successful frontier-scale models and cross-domain adoption could make this a foundational platform.

Proposed pre-seed positioning: approximately US\$10M priced pre-money, conditional on team, rights, budget and investor diligence. Capital should buy the next decisive model, runtime and commercial proof points. A larger valuation must be earned through reproducible advantage and adoption.

Prepared 3 October 2026. [Full investment case](#) and [evidence manifest](#). This is a fundraising thesis; commercial success, clockless chip energy savings and frontier supremacy remain to be demonstrated.